

```
from google.colab import files
uploaded = files.upload()
```

[Choose Files](#) Invistico_Airline.csv

Invistico_Airline.csv(text/csv) - 11960454 bytes, last modified: 6/25/2026 - 100% done
Saving Invistico_Airline.csv to Invistico_Airline.csv

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, precision_score, recall_score

# Load the dataset
df = pd.read_csv("Invistico_Airline.csv")

# Display the first 5 rows and target variable distribution
print(df.head())
print("\nTarget Variable Distribution:")
print(df['satisfaction'].value_counts())
```

```
satisfaction  Customer Type  Age  Type of Travel  Class \
0  satisfied  Loyal Customer  65  Personal Travel  Eco
1  satisfied  Loyal Customer  47  Personal Travel  Business
2  satisfied  Loyal Customer  15  Personal Travel  Eco
3  satisfied  Loyal Customer  60  Personal Travel  Eco
4  satisfied  Loyal Customer  70  Personal Travel  Eco

Flight Distance  Seat comfort  Departure/Arrival time convenient \
0      265      0      0
1     2464      0      0
2     2138      0      0
3      623      0      0
4      354      0      0

Food and drink  Gate location  ...  Online support  Ease of Online booking \
0      0      2  ...      2      3
1      0      3  ...      2      3
2      0      3  ...      2      2
3      0      3  ...      3      1
4      0      3  ...      4      2

On-board service  Leg room service  Baggage handling  Checkin service \
0      3      0      3      5
1      4      4      4      2
2      3      3      4      4
3      1      0      1      4
4      2      0      2      4

Cleanliness  Online boarding  Departure Delay in Minutes \
0      3      2      0
1      3      2      310
2      4      2      0
3      1      3      0
4      2      5      0

Arrival Delay in Minutes
0      0.0
1     305.0
2      0.0
3      0.0
4      0.0

[5 rows x 22 columns]

Target Variable Distribution:
satisfaction
satisfied      71087
dissatisfied    58793
Name: count, dtype: int64
```

```
# Check for missing values
print("Missing values per column:")
print(df.isnull().sum())

# Drop rows with missing values (if any, e.g., in Arrival Delay)
```

```

df = df.dropna()

# Identify categorical columns (excluding the target 'satisfaction')
categorical_cols = df.select_dtypes(include=['object']).columns.drop('satisfaction')

# One-hot encode categorical predictors
df_encoded = pd.get_dummies(df, columns=categorical_cols, drop_first=True)

# Binary encode the target variable 'satisfaction' (1 for Satisfied, 0 for Dissatisfied)
# Adjust the string mapping if your dataset uses 'satisfied'/'dissatisfied' lowercase
df_encoded['satisfaction'] = df_encoded['satisfaction'].map({'satisfied': 1, 'satisfied': 1, 'Satisfied': 1, 'neutral or dissat

# Display the new columns to verify encoding
print("\nEncoded Dataset Columns:")
print(df_encoded.columns.tolist()[:10]) # showing first 10 columns

```

```

Missing values per column:
satisfaction                0
Customer Type               0
Age                         0
Type of Travel              0
Class                      0
Flight Distance             0
Seat comfort                0
Departure/Arrival time convenient 0
Food and drink              0
Gate location               0
Inflight wifi service       0
Inflight entertainment      0
Online support              0
Ease of Online booking      0
On-board service            0
Leg room service            0
Baggage handling            0
Checkin service             0
Cleanliness                 0
Online boarding             0
Departure Delay in Minutes  0
Arrival Delay in Minutes   393
dtype: int64

Encoded Dataset Columns:
['satisfaction', 'Age', 'Flight Distance', 'Seat comfort', 'Departure/Arrival time convenient', 'Food and drink', 'Gate location

```

```

# Separate features (X) and target variable (y)
X = df_encoded.drop(columns=['satisfaction'])
y = df_encoded['satisfaction']

# Split the dataset (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

print("Training set shape:")
print(f"X_train shape: {X_train.shape}, y_train shape: {y_train.shape}")
print("\nTesting set shape:")
print(f"X_test shape: {X_test.shape}, y_test shape: {y_test.shape}")

```

```

Training set shape:
X_train shape: (103589, 22), y_train shape: (103589,)

Testing set shape:
X_test shape: (25898, 22), y_test shape: (25898,)

```

```

from sklearn.preprocessing import StandardScaler

# Identify continuous numerical columns to scale
# We look for columns that aren't binary (0 or 1) dummy variables
numerical_cols = X_train.select_dtypes(include=['int64', 'float64']).columns

# Initialize the scaler
scaler = StandardScaler()

# Fit the scaler on the training data and transform both train and test sets
# This centers the data around 0 with a standard deviation of 1
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

X_train_scaled[numerical_cols] = scaler.fit_transform(X_train[numerical_cols])
X_test_scaled[numerical_cols] = scaler.transform(X_test[numerical_cols])

```

```
# Re-initialize and train the Logistic Regression model on the scaled data
model = LogisticRegression(max_iter=1000, random_state=42)
model.fit(X_train_scaled, y_train)

# Make predictions using the scaled features
y_pred = model.predict(X_test_scaled)
y_pred_proba = model.predict_proba(X_test_scaled)[: , 1]

print("Model training complete with scaled features!")
```

Model training complete with scaled features!

```
# Generate the Confusion Matrix
cm = confusion_matrix(y_test, y_pred)


# Calculate key evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print("--- Model Evaluation Metrics ---")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print("\nConfusion Matrix:")
print(cm)

# Optional: Visualize the confusion matrix cleanly using Seaborn
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Dissatisfied', 'Satisfied'],
            yticklabels=['Dissatisfied', 'Satisfied'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix Breakdown')
plt.show()
```

```
--- Model Evaluation Metrics ---
Accuracy: 0.8253
Precision: 0.8393
Recall: 0.8420
```

```
Confusion Matrix:
[[ 9436 2285]
 [ 2240 11937]]
```



```
# Create a DataFrame to view the coefficients of each feature
coefficients_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_[0]
})

# Calculate the exponential of the coefficients to get the Odds Ratios
coefficients_df['Odds Ratio'] = np.exp(coefficients_df['Coefficient'])
```

```
# Sort by the absolute value of the coefficient to see the strongest drivers first
coefficients_df['Abs_Coefficient'] = coefficients_df['Coefficient'].abs()
coefficients_df = coefficients_df.sort_values(by='Abs_Coefficient', ascending=False).drop(columns=['Abs_Coefficient'])

print("--- Model Coefficients and Odds Ratios ---")
print(coefficients_df)
```

```
--- Model Coefficients and Odds Ratios ---
```

	Feature	Coefficient	Odds Ratio
18	Customer Type_disloyal Customer	-1.890566	0.150986
7	Inflight entertainment	0.973634	2.647548
21	Class_Eco Plus	-0.775057	0.460678
19	Type of Travel_Personal Travel	-0.758093	0.468559
20	Class_Eco	-0.725599	0.484035
2	Seat comfort	0.408756	1.504945
10	On-board service	0.399074	1.490444
13	Checkin service	0.357956	1.430403
3	Departure/Arrival time convenient	-0.332580	0.717071
9	Ease of Online booking	0.315901	1.371494
4	Food and drink	-0.313184	0.731115
11	Leg room service	0.310542	1.364164
17	Arrival Delay in Minutes	-0.304175	0.737731
15	Online boarding	0.186024	1.204451
1	Flight Distance	-0.182215	0.833422
5	Gate location	0.170317	1.185680
8	Online support	0.151236	1.163271
0	Age	-0.130895	0.877310
6	Inflight wifi service	-0.120173	0.886767
12	Baggage handling	0.109432	1.115644
16	Departure Delay in Minutes	0.107958	1.114001
14	Cleanliness	0.067550	1.069884

```
from sklearn.metrics import classification_report

# Generate the complete classification report
report = classification_report(y_test, y_pred, target_names=['Dissatisfied', 'Satisfied'])

print("--- Scikit-Learn Classification Report ---")
print(report)
```

```
--- Scikit-Learn Classification Report ---
```

	precision	recall	f1-score	support
Dissatisfied	0.81	0.81	0.81	11721
Satisfied	0.84	0.84	0.84	14177
accuracy			0.83	25898
macro avg	0.82	0.82	0.82	25898
weighted avg	0.83	0.83	0.83	25898


```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(9, 5))

# Plotting the distribution of satisfaction across different entertainment ratings
sns.countplot(data=df, x='Inflight entertainment', hue='satisfaction', palette='Blues')

plt.title('Passenger Satisfaction Distribution by In-flight Entertainment Rating', fontsize=14, pad=15)
plt.xlabel('In-flight Entertainment Rating (1 - 5)', fontsize=12)
plt.ylabel('Number of Passengers', fontsize=12)
plt.legend(title='Customer Status', labels=['Dissatisfied', 'Satisfied'])
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```



Start coding or [generate](#) with AI.